

Introduction to R

Guy J. Abel

R Console Start Up

```
R version 4.0.2 (2020-06-22) -- "Taking Off Again"
Copyright (C) 2020 The R Foundation for Statistical Computing
Platform: x86_64-w64-mingw32/x64 (64-bit)
```

R is free software and comes with ABSOLUTELY NO WARRANTY.
You are welcome to redistribute it under certain conditions.
Type 'license()' or 'licence()' for distribution details.

R is a collaborative project with many contributors.
Type 'contributors()' for more information and
'citation()' on how to cite R or R packages in publications.

Type 'demo()' for some demos, 'help()' for on-line help, or
'help.start()' for an HTML browser interface to help.

Type 'q()' to quit R.

>

R Console Start Up

- Note the > at the bottom.
- Whenever you see this symbol, it means that R is not doing anything and just waiting for your input.
- It's called the prompt.

Very Basics

- We can enter R commands at the prompt

```
> 3 - 1  
[1] 2
```

- When R is running it stores everything in the computer's active memory.
 - R knows 1 and 3, as well as the basic operator `-`.
 - It creates an object in the active memory containing the results of the computation.
 - R prints the content of the results object to the screen.
- R let's you know it has finished when displaying in the square brackets `[]`
 - The square brackets `[]` tells you how many results were calculated.
 - In our case this object contains only one value, i.e. R has finished and printed the first and only value `[1]` of the computation.
- You can scroll through previous commands you've entered by using the up ("`→`") and down ("`↓`") keys on your keyboard.

R Basics

- R may sometimes display a + (also known as the continuation line).
 - R will not report the [], as it is not displaying any content from an object stored in memory.
 - Entered an incomplete command
 - R is waiting for more input.

```
> 3 -  
+ 
```

- If you are not sure what's going on hit Esc or Ctrl-C.
 - It will tell R to forget it and bring back the prompt.

R Basics

- If you know what R wants, then you can complete the computation:

```
> 3 -
+ 1
[1] 2
```

- The continuation line also commonly occurs with
 - Un-closed parenthesis ()
 - Un-closed quotes ""

```
> 7 / (1 + 3
+ )
[1] 1.75
> "Guy
+ "
[1] "Guy\n"
```

- Remember if you are in a syntactic hole... Esc.

Calculator

- R understands the following basic operators:
 - ① + and - for addition and subtraction
 - ② * and / for multiplication and division
 - ③ ^ for raising to the power (e.g. $10^2 = 100$)
- R observes standard rules of operator precedence.
- You can use brackets () if you are not sure, e.g.

```
> 7 / (1 + 3)
[1] 1.75
```

is not the same as this:

```
> 7 / 1 + 3
[1] 10
```

Character Strings

- R also allows you to type in strings of characters (letters, words, phrases)
- If you do not use quotation marks R will think you are asking for an object

```
> "Guy"  
[1] "Guy"  
> 'Guy'  
[1] "Guy"
```

```
> Guy  
Error: object 'Guy' not found
```

- We will go into depth on objects a little later.
- For now remember text must go in quotation.

Comments

- The comment operator # will tell R to ignore everything printed after it (in the current line).
- Extremely useful to annotate your code.

```
> 1 + 2 + 3 # some additions  
[1] 6
```

- It is good practice to annotate your code.
 - Within a surprisingly short period of time you will forget what each bit of code is trying to do.

Comments

- Be careful.

```
> 3 - 1 # + 20  
[1] 2
```

- Misplaced comments can break your code.

```
> 3 - # 1  
+
```

Spacing

- For the most part, R does not care about spacing.

```
> 3 - 1
[1] 2
>
> # same as
> 3      -          1
[1] 2
```

- For character strings spaces matter...

```
> message("Strings obey spacing.")
Strings obey spacing.
> message(" Strings obey      spacing .  ")
Strings obey      spacing .
```

Semi-Colons

- R evaluates code line by line.
- A line break tells R that a statement is to be evaluated.
- Instead of a line break, you can use a semicolon (;) to tell R where statements end.

```
> "Guy"  
[1] "Guy"
```

```
> 3 - 1    5 * 9  
Error: unexpected numeric constant in " 1 + 2    5"
```

```
> 3 - 1; 5 * 9  
[1] 2  
[1] 45
```

Exercise 1 (ex11.R)

- Open ex11.R and complete the following exercises:

```
# 1. 5 plus 6
```

```
# 2. 2 multiplied by 8
```

```
# 3. 8 divided by 3
```

```
# 4. 909 minus 506
```

```
# 5. 5 to the power of 10
```

```
# 6. Tell R to say your first name
```

```
# 7. Q: Which symbol does R use to ignore all code after?
```

```
#    A:
```

```
# 8. On one line of code with two R print commands write
```

```
#    your first and last name.
```

Basic Functions

- R comes with a many many pre-installed functions.
 - Installed as part of the base package which is located in your library directory.
 - Functions have names and take arguments in parentheses: `function(...)`

- Takes the form

```
function(argument1, argument2)
```

- Arguments are options for the function.
- Each function has different arguments.
- If there are multiple arguments, use a comma , to separate.

Basic Functions

- We can find out what arguments (options) are for a function using ?

```
> ?log
```

- In RStudio you can also use Tab once you have typed the function name and opened the parenthesis
- The help file reports that the function takes two arguments, separated by a comma.
 - x a number you want to take the logarithm of (no default)
 - base the base system for the logarithm (default base = exp(1))

```
> log(x = 10)
[1] 2.302585
> # same as
> log(x = 10, base = exp(1))
[1] 2.302585
> # change the base argument...
> log(x = 10, base = 10)
[1] 1
```

Basic Functions

- R knows what arguments are supplied to the function and their order.

```
> log(x = 10)
[1] 2.302585
> # same as
> log(10)
[1] 2.302585
>
> # knows the order as well (i.e. second input is the `base`)
> log(10, 10)
[1] 1
> # same as
> log(x = 10, base = 10)
[1] 1
```

- Whilst this reduces your typing, I advise not to do it.
 - Difficult to remember which inputs for which arguments.
 - Even more difficult for others looking at your code.
 - Using Tab and Enter in RStudio means its easy to write argument names fully.

Basic Functions

Basic mathematical functions in base R

Function	Description
<code>log()</code>	computes natural logarithms
<code>log10()</code>	computes logarithm to the base 10
<code>exp()</code>	computes the exponential function
<code>sqrt()</code>	takes the square root
<code>abs()</code>	returns the absolute value
<code>factorial()</code>	returns the factorial
<code>sign()</code>	returns the sign (negative or positive)
<code>round()</code>	rounds the input to the desired digit

Remember we can use `?` to view the help file of any function, e.g. `?round`

Exercise 2 (ex12.R)

1. Natural logarithm of 5

2. Square root of 121

3. Absolute value of 10 and -11

4. $8 \times 7 \times 6 \times 5 \times 4 \times 3 \times 2 \times 1$ (the factorial of 8)

5. Round pi (3.141593) to 3 decimal places

6. The logarithm of 100 to the base 10

7. The exponential of 4

8. Print your name using the message function

Logic Operators

- Among the most used features of R are logical operators.
- You will use these throughout your code and they are crucial for all sorts of data manipulation.
- When R evaluates statements containing logical operators it will return either TRUE or FALSE

Function	Description
<	less than
<=	less than or equal to
>	greater than
>=	greater than or equal to
==	equal
!=	not equal
&	and
	or

Logic Operators

```
> pi
[1] 3.141593
>
> pi > 5
[1] FALSE
>
> pi > 3 & pi > 4
[1] FALSE
>
> pi > 5 | pi < 4
[1] TRUE
```

Exercise 3 (ex13.R)

1. Test if 3 is larger than pi

2. Test if pi is equal to 3.141593

3. Test if 7 divided by 3 is less than or equal to 3

4. Test if 5 times 2 is greater than or equal to 10

5. Test if 1 plus 5 is not equal to 7

6. Test if logarithm of 1000 is larger 7

7. Test if pi is greater than 3 and 4

8. Test if pi is less than 3 or 4

Assignment and Reference

- Instead of recalculating everything over and over again we can give things names and recall them later.
- Before we implicitly relied on and then manipulated objects and R implicitly printed these objects to the screen.

```
> 1 + 2  
[1] 3
```

- We can assign and recall names using the assignment operator `<-`.
 - Think of this as the `M+` button on your calculator.

```
> a <- 1 + 2
```

- Above, R no longer gave the answer to our problem. It just returns the prompt.

Assignment and Reference

- We can print the results to the screen by typing the name of our new object

```
> a  
[1] 3
```

- It does not do the calculation again when you print the object.

Beware

- We can use just about any name we like except
 - Cannot be a number, e.g. `3 <- 1 + 2`
 - Cannot start with with a number, e.g. `3a <- 1 + 2`
- You can break R code by creating objects that already are existing functions.

```
> exp
function (x)  .Primitive("exp")
> exp <- 1 + 3
> exp
[1] 4
```

- The exp function still works, but is lower down in the search environment than the new object.
 - We could create our own function exp that does something completely different.
 - Directly get a function in a specific package using `::`

```
> base::exp
function (x)  .Primitive("exp")
> readxl::read_excel
function (path, sheet = NULL, range = NULL, col_names = TRUE,
         col_types = NULL, na = "", trim_ws = TRUE, skip = 0, n_max = Inf,
         guess_max = min(1000, n_max), progress = readxl_progress(),
         .name_repair = "unique")
```


Beware

- R is case-sensitive.

```
> A
Error: object 'A' not found
```

```
> a
[1] 3
```

- If you are not sure, check the potential object name by printing it before assigning it!

```
> e
Error in eval(expr, envir, enclos): object 'e' not found
```

Playing with Trivial Vectors

- Named objects behave just like the ones R already knows.
- This is very useful when things get more complicated:

```
> a
[1] 3
> a * 2
[1] 6
>
> b <- a * sqrt(a)
> b
[1] 5.196152
```

Keeping Track of Objects

- To see what objects you have created (the ones R stored in active memory) you can use the `ls()` function.

```
> ls()
[1] "a"    "b"    "exp"
```

- RStudio also shows each created object in the Environment tab.
- If you want to remove an object from memory use the `rm()` function.
- Be very careful. This will delete thing permanently.

```
> rm(b, exp)
> ls()
[1] "a"
```

- If you want to remove all objects from active memory can also use the button with a brush on it above the object list in RStudio or directly in the console:

```
> rm(list = ls())
```

Real Vectors

- So far we have only created trivial vectors of length 1. Let's assign some longer ones.
- To do this you will use the `c()` function.
 - The `c` stands for concatenate,
 - Combines a bunch of elements together, separated by commas.
 - Make sure you never call an object `c` (like we just created objects `a` and `b`)

```
> v1 <- c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
> v1
[1] 1 2 3 4 5 6 7 8 9 10
```

How about a character vector?

```
> v2 <- c("a", "b", "c", "d")
> v2
[1] "a" "b" "c" "d"
```

Or ...

```
> # will convert everything to characters if there is at least one
> v3 <- c(1, "two", 3, 4)
> v3
[1] "1" "two" "3" "4"
```

Real Vectors

- You can also string multiple vectors together with the `c()` function.

```
> v4 <- c(v1 , v2 , v1)
> v4
 [1] "1"  "2"  "3"  "4"  "5"  "6"  "7"  "8"  "9"  "10" "a"  "b"  "c"  "d"  "1"
[16] "2"  "3"  "4"  "5"  "6"  "7"  "8"  "9"  "10"
```

Simplifying Vector Creation

- Using the `c()` function can be tedious
 - Have to manually type all elements of a vector.
 - Will make mistakes
- Use the colon `(:)` to tell R to create an integer vector.

```
> 1:20
[1] 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20
> # backwards
> 10:-5
[1] 10 9 8 7 6 5 4 3 2 1 0 -1 -2 -3 -4 -5
```

- Use the `seq()` function for more general sequences.

```
> seq(from = 0, to = 10)
[1] 0 1 2 3 4 5 6 7 8 9 10
> # the 'by' argument let's you set the increments
> seq(from = 0, to = 10, by = 2)
[1] 0 2 4 6 8 10
```

Simplifying Vector Creation

```
> # the length.out argument specifies the
> # length of the vector and R figures out
> # the increments itself
> seq(from = 0, to = 10, length.out = 25)
 [1] 0.0000000 0.4166667 0.8333333 1.2500000 1.6666667 2.0833333
 [7] 2.5000000 2.9166667 3.3333333 3.7500000 4.1666667 4.5833333
[13] 5.0000000 5.4166667 5.8333333 6.2500000 6.6666667 7.0833333
[19] 7.5000000 7.9166667 8.3333333 8.7500000 9.1666667 9.5833333
[25] 10.0000000
```

Simplifying Vector Creation

- Another useful function is `rep()` which allows you to repeat things.

```
> rep(x = 0, times = 10)
[1] 0 0 0 0 0 0 0 0 0 0
> # as always you can drop the argument name
> rep(x = "Hello", times = 3)
[1] "Hello" "Hello" "Hello"
> # repeating vector 1 twice
> rep(x = v1 , times = 2)
[1] 1 2 3 4 5 6 7 8 9 10 1 2 3 4 5 6 7 8 9 10
> # we can repeat each element as well
> rep(x = v2 , each = 2)
[1] "a" "a" "b" "b" "c" "c" "d" "d"
```


Vector Operations

- R is very powerful when working with vectors.
- Most standard mathematical functions work with vectors.

```
> v1
[1] 1 2 3 4 5 6 7 8 9 10
> v1 + v1
[1] 2 4 6 8 10 12 14 16 18 20
> v1 ^ 2
[1] 1 4 9 16 25 36 49 64 81 100
> log(x = v1)
[1] 0.0000000 0.6931472 1.0986123 1.3862944 1.6094379 1.7917595 1.9459101
[8] 2.0794415 2.1972246 2.3025851
```

Exercise 4 (ex14.R)

1. Create a vector called a1 for the 1, 4, 9, 16, 25

2. Create a vector called a2 from 1 to 100

3. Create a vector called a3 from -8 to 20 in steps of 4

4. Create a vector called a4 that combines a1 and a3

5. Create a vector called a5 that has your first and last name repeated 3 times

6. Find the square root of each element in a1

7. Test to see where a3 are positive (greater than zero)

8. Divide a2 by a1. Then divide a2 by a3. What is going on?

Selected Functions for Vectors

- R has many functions in the base package that work with vectors.

Function	Description
<code>sum()</code>	sums of the elements of the vector
<code>prod()</code>	product of the elements of the vector
<code>min()</code>	minimum of the elements of the vector
<code>max()</code>	maximum of the elements of the vector
<code>range()</code>	the range of the vector
<code>sort()</code>	sorts the vector (argument: <code>decreasing = FALSE</code>)
<code>rev()</code>	Reverse the order of the vector
<code>length()</code>	returns the length of the vector
<code>which()</code>	returns the index after evaluating a logical statement
<code>unique()</code>	returns a vector of all the unique elements of the input
<code>table()</code>	returns tabulation count for each unique element
<code>ifelse()</code>	returns different values depending on logical test

Conditionals

- The `ifelse()` function is used when you want one action when statement TRUE and different action when FALSE.
- Very useful when creating variables in data frames. Has three arguments
 - test the statement
 - yes action if statement is TRUE
 - no action if statement is FALSE

```
> v1
[1] 1 2 3 4 5 6 7 8 9 10
>
> # logical test
> v1 > 5
[1] FALSE FALSE FALSE FALSE FALSE TRUE TRUE TRUE TRUE TRUE
>
> # display a character vector for results
> ifelse(test = v1 > 8, yes = "High", no = "Low")
[1] "Low" "Low" "Low" "Low" "Low" "Low" "Low" "Low" "High" "High"
>
> # can return numeric values as well
> ifelse(test = v1 > 8, yes = 1, no = 0)
[1] 0 0 0 0 0 0 0 0 1 1
```

Exercise 5 (ex15.R)

1. Find the sum of all elements in a1

2. What is the largest element in a4

3. Multiply all elements of a2 together

4. How many elements are there in a5

5. Reverse the order of a2

6. Which elements of a1 are greater than 10

7. What is the range of values in a4

8. What are the unique elements of a4

9. Sort the values in a4

10. Create a tabulation of the elements in a4

11. Create a vector from a4 which recodes elements above zero as 1 and below zero

12. Create a vector from a4 which recodes elements above zero with label "positive"
and below zero or zero with label "negative or zero"

Selected Statistical Functions for Vectors

- R has many statistical functions in the stats package that work with vectors.

Function	Description
<code>mean()</code>	mean of the elements
<code>median()</code>	median of the elements
<code>sd()</code>	the standard deviation
<code>var()</code>	the variance
<code>cov()</code>	the co-variance (takes two inputs <code>cov(x,y)</code>)
<code>cor()</code>	the correlation coefficient (takes two inputs <code>cor(x,y)</code>)
<code>IQR()</code>	Inter Quartile Range
<code>quantile()</code>	returns values corresponding to the given probabilities (takes input probs)
<code>summary()</code>	returns summary statistics

Sampling

- R can generate many different types of random numbers.
- The `sample()` function draws randomly from a given vector

```
> v1
[1] 1 2 3 4 5 6 7 8 9 10
> sample(x = v1, size = 5)
[1] 1 7 3 8 5
> sample(x = v1, size = 15, replace = TRUE)
[1] 10 6 10 5 3 10 3 8 3 6 1 2 8 7 8
```

Sampling

- Sample from many different types of statistical distributions

Function	Description
<code>rnorm()</code>	random generation for the normal distribution (<code>n</code> , <code>mean</code> , <code>sd</code>)
<code>rbinom()</code>	random generation for the binomial distribution (<code>n</code> , <code>size</code> , <code>prob</code>)
<code>rpois()</code>	random generation for the Poisson distribution (<code>n</code> , <code>lambda</code>)
<code>runif()</code>	random generation for the uniform distribution (<code>n</code> , <code>min</code> , <code>max</code>)

```
> # normal
> rnorm(n = 5, mean = 1, sd = 2)
[1] -1.5948555 2.4919915 0.8964033 3.5129072 -1.1273850
> # binomial
> rbinom(n = 8, size = 10, prob = 0.3)
[1] 3 3 1 2 2 2 2 4
> # poisson
> rpois(n = 8, lambda = 2)
[1] 2 0 1 0 0 2 2 0
> # uniform
> runif(n = 5, min = 5, max = 8)
[1] 7.875833 6.655136 7.239813 7.612478 6.281276
```


Exercise 6 (ex16.R)

1. Create a vector a6 that is a sequence from 1 to 25 of length 8

a6

2. What is the mean of a6

3. what is the standard deviation of a6

*# 4. Create a vector a7 that is a random sample of size 8 from
sequence of integers between 1 to 100*

a7

5. What is the correlation of a6 and a7

*# 6. Create a vector a8 of length 8 of random numbers from a normal
distribution with mean 10 and standard deviation 2*

7. What is Inter Quartile Range of a8

8. Create a summary of a8

Objects

- R is an object oriented language.
 - Everything in R is an object.
 - When R does anything, it creates and manipulates objects.
 - Objects are a bit like the memory button on the calculator, but with much much more freedom.
- R objects come in different types and flavors.

Basic Objects: Vectors

- One-dimensional sequences of elements of the same mode. (More on modes later)
- For example, this could be vector of length 26 (i.e. one containing 26 elements) where each element is a letter in the alphabet.
- R has some built in vectors

```
> pi
[1] 3.141593
> LETTERS
[1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q" "R" "S"
[20] "T" "U" "V" "W" "X" "Y" "Z"
```

Basic Objects: Matrices

- Two dimensional rectangular objects (matrices) or
- All elements of matrices have to be of the same mode.
- Can build matrices using the `matrix()` function

```
> matrix(data = 1:12, nrow = 4)
```

```
  [,1] [,2] [,3]
```

```
[1,]    1    5    9
```

```
[2,]    2    6   10
```

```
[3,]    3    7   11
```

```
[4,]    4    8   12
```

Basic Objects: Data Frames

- Data frames are best understood as special matrices (technically they are a type of list).
- They are two dimensional containers with
 - Rows corresponding to observations
 - Columns corresponding to vectors
- R has some built in data frames

```
> swiss
```

	Fertility	Agriculture	Examination	Education	Catholic
Courtelary	80.2	17.0	15	12	9.96
Delemont	83.1	45.1	6	9	84.84
Franches-Mnt	92.5	39.7	5	5	93.40
Moutier	85.8	36.5	12	7	33.77
Neuveville	76.9	43.5	17	15	5.16
Porrentruy	76.1	35.3	9	7	90.57
Broye	83.8	70.2	16	7	92.85
Glane	92.4	67.8	14	8	97.16
Gruyere	82.4	53.3	12	7	97.67
Sarine	82.9	45.2	16	13	91.38
Veveyse	87.1	64.5	14	6	98.61
Aigle	64.1	62.0	21	12	8.52
Aubonne	66.9	67.5	14	7	2.27
Avenches	68.9	60.7	19	12	4.43
Cossonay	61.7	69.3	22	5	2.82

Basic Objects: Lists

- Like vectors but they do not have to contain elements of the same mode.
- The first element of a list could be a vector of the 26 letters of the alphabet.
- The second element could contain a vector of numbers.
- A third could be a 2 by 3 matrix.

```
> list(LETTERS, 1:10, matrix(1:6, nrow = 2))
[[1]]
 [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O" "P" "Q" "R" "S"
[20] "T" "U" "V" "W" "X" "Y" "Z"

[[2]]
 [1] 1 2 3 4 5 6 7 8 9 10

[[3]]
      [,1] [,2] [,3]
[1,]    1    3    5
[2,]    2    4    6
```

Basic Objects: Lists

- Lists are used in R for regression results and many other complex outputs:
- We can check the structure of objects using the `str()` function:

```
> m <- lm(formula = Fertility ~ Education + Catholic, data = swiss)
```

```
> str(m)
```

List of 12

```
$ coefficients : Named num [1:3] 74.234 -0.788 0.111
..- attr(*, "names")= chr [1:3] "(Intercept)" "Education" "Catholic"
$ residuals : Named num [1:47] 14.32 6.55 11.85 13.34 13.92 ...
..- attr(*, "names")= chr [1:47] "Courtelary" "Delemont" "Franches-Mnt" "Moutier"
$ effects : Named num [1:47] -480.9 -56.2 -31 11.5 13.2 ...
..- attr(*, "names")= chr [1:47] "(Intercept)" "Education" "Catholic" "" ...
$ rank : int 3
$ fitted.values: Named num [1:47] 65.9 76.5 80.7 72.5 63 ...
..- attr(*, "names")= chr [1:47] "Courtelary" "Delemont" "Franches-Mnt" "Moutier"
$ assign : int [1:3] 0 1 2
$ qr :List of 5
..$ qr : num [1:47, 1:3] -6.856 0.146 0.146 0.146 0.146 ...
.. ..- attr(*, "dimnames")=List of 2
.. ..$ : chr [1:47] "Courtelary" "Delemont" "Franches-Mnt" "Moutier" ...
.. ..$ : chr [1:3] "(Intercept)" "Education" "Catholic"
.. ..- attr(*, "assign")= int [1:3] 0 1 2
..$ qraux: num [1:3] 1.15 1.03 1.17
..$ pivot: int [1:3] 1 2 3
..$ tol : num 1e-07
```

Modes

- All objects have a certain mode.
- Some objects can only deal with one mode at a time, others can store elements of multiple modes.
- Some basic modes include:
 - 1 integer: integers (e.g. 1, 2 or -1000)
 - 2 numeric: real numbers (e.g 2.336, -0.35)
 - 3 character: elements made up of text-strings (e.g. "text", "Hello World!", or "123")
 - 4 logical: data containing logical constants (i.e. TRUE and FALSE)

Indexing

- Sometimes you do not want to print or manipulate an entire vector.
- This is where indexing comes in - helps select or adjust values of your R object based on **position** integer(s)
- Indexing vectors is done with `[]`.

```
> v4 <- c("I", "really", "like", "chocolate", "ice cream.")
> # with the bracket we reference the third element
> v4[3]
[1] "like"
> # we can reference a sequence of elements
> v4[2:4]
[1] "really"      "like"        "chocolate"
> # or any elements we like
> v4[c(1,3,4)]
[1] "I"           "like"        "chocolate"
```

Indexing

```
> # all except the 2nd element
> v4[-2]
[1] "I"           "like"         "chocolate"   "ice cream."
> # and we can change elements
> v4[4] <- "strawberry"
> v4
[1] "I"           "really"       "like"         "strawberry"  "ice cream."
```

Indexing

- Indexing also works with matrices
- Using a , to separate rows and columns [row, column]

```

> # create a matrix
> m0 <- matrix(data = 1:12, nrow = 3)
> m0

      [,1] [,2] [,3] [,4]
[1,]    1    4    7   10
[2,]    2    5    8   11
[3,]    3    6    9   12
> # bottom right element
> m0[3, 4]
[1] 12
> # third column
> m0[, 3]
[1] 7 8 9
  
```

Indexing

- The [row, column] indexing also works with data frames
- R users usually prefer to use the name of the column
- The dataframe\$column selects the entire column

```
> # can use the raw row and column coordinates
```

```
> swiss[38,1]
```

```
[1] 79.3
```

```
> # the row and column names
```

```
> swiss["Sion","Fertility"]
```

```
[1] 79.3
```

```
> # the column name and row coordinates
```

```
> # this is useful if there are no row names, which can often be the case.
```

```
> swiss$Fertility[38]
```

```
[1] 79.3
```

```
> # the whole column
```

```
> swiss$Fertility
```

```
[1] 80.2 83.1 92.5 85.8 76.9 76.1 83.8 92.4 82.4 82.9 87.1 64.1 66.9 68.9 61.7
[16] 68.3 71.7 55.7 54.3 65.1 65.5 65.0 56.6 57.4 72.5 74.2 72.0 60.5 58.3 65.4
[31] 75.5 69.3 77.3 70.5 79.4 65.0 92.2 79.3 70.4 65.7 72.7 64.4 77.6 67.6 35.0
[46] 44.7 42.8
```

Indexing

- Indexing also works for lists
- Can use the \$ or two square brackets:

```
> m$coefficients
(Intercept)    Education    Catholic
  74.2336892   -0.7883293    0.1109210
> m[["coefficients"]]
(Intercept)    Education    Catholic
  74.2336892   -0.7883293    0.1109210
```

Missing Values

- NA is used to represent missing data
- Different functions treat missing data differently

```
> v5 <- c(10, 7, NA, NA, 0, NA, -2, 8)
> v5
[1] 10 7 NA NA 0 NA -2 8
> # this won't work by default
> max(v5)
[1] NA
> # there is usually an option to allow functions to omit NA
> max(v5, na.rm = TRUE)
[1] 10
> # test for missing values using is.na
> is.na(v5)
[1] FALSE FALSE TRUE TRUE FALSE TRUE FALSE FALSE
> # test for non missing values using !is.na
> !is.na(v5)
[1] TRUE TRUE FALSE FALSE TRUE FALSE TRUE TRUE
> # this works (does not have na.rm argument)
> summary(v5)
  Min. 1st Qu.  Median    Mean 3rd Qu.    Max.   NA's
 -2.0     0.0     7.0     4.6     8.0    10.0     3
```

Exercise 7 (ex17.R)

```
# 1. Print out the second row of the m1 matrix, where
# m1 <- matrix(1:28, nrow = 4, ncol = 7)

# 2. Print out the top right value of the m1 matrix

# 3. Print out the first column of the m1 matrix

# 4. Print out the whole of the Infant.Mortality column in the swiss data

# 5. What is the maximum fertility in the swiss data

# 6. Print out the Catholic value for Glane in the swiss data

# 7. Set the third and sixth element of a8 to missing values

# 8. What is the variance of the remaining numbers in a8
```

Functions

- R is a programming language. We can create (program) our own functions using the `function()` function.
- Wrap up code for repetitive tasks
- Simplifies complex tasks
- Creates user friendly access to your R code.
- Functions takes the format

```
function(arguments, ...){
  actions
  return(value)
}
```

- Wrap all sorts of actions inside of a function to simplify repetitive tasks.
- Specify as many arguments as you like.
- Set argument(s) to `NULL` if we do not want a default value. Forces the user to set.
- The `return()` function is not essential.
- The actions will take place but nothing is outputted unless we print the object on the last line of the function.

Functions

```
> my_square <- function(x = NULL){  
+   y <- x ^ 2  
+   return(y)  
+ }  
> my_square(x = 2)  
[1] 4  
> my_square(x = 10)  
[1] 100  
>  
> # add a p argument to raise x to any power  
> my_power <- function(x = NULL, p = 2){  
+   y <- x ^ p  
+   return(y)  
+ }  
> my_power(x = 2)  
[1] 4  
> my_power(x = 2, p = 4)  
[1] 16
```

Functions

```
> my_square <- function(x = NULL){  
+   y <- x ^ 2  
+   return(y)  
+ }  
> my_square(x = 2)  
[1] 4  
> my_square(x = 10)  
[1] 100  
>  
> # add a p argument to raise x to any power  
> my_power <- function(x = NULL, p = 2){  
+   y <- x ^ p  
+   return(y)  
+ }  
> my_power(x = 2)  
[1] 4  
> my_power(x = 2, p = 4)  
[1] 16
```

Loops

- Loops are functions that carry out repetitive actions.
- The `for()` function in R is used to repeat actions a given number of times.

General format:

```
for(element in vector){
  action
}
```

- Within the `{}` we put the actions we want repeated multiple times.
- General code, using `element` in place of the parts of the code that are non-general.
- Within the `()` we put an instruction.
- The `element` values are taken from `vector` in each round of the loop.
- Typically people use a single letter such as `i`.
- The vector contains all the possible `element` values to be used.

Loops

```
> # create empty vector to save action
> v8 <- rep(x = NA, times = 12)
> v8
[1] NA NA NA NA NA NA NA NA NA NA NA NA
>
> # fill the vector with random numbers from a poisson distribution
> # with increasing lambda mean based on i value
> for(i in 1:10){
+   x <- rpois(n = 1, lambda = i)
+   message(x)
+   v8[i] <- x
+ }
4
0
4
3
8
1
5
5
9
8
> v8
[1] 4 0 4 3 8 1 5 5 9 8 NA NA
```

Loops

- Loops can be really useful for repetitive tasks

```
> for(i in 1:10){  
+   # read in different excel sheets  
+   d <- read_excel(path = "myexcelfile.xlsx", sheet = i)  
+   # run regression models using data from excel sheet  
+   m <- lm(formula = y ~ x1 + x2, data = d)  
+   # save regression model coefficients  
+   write_csv(x = m$coefficients, path = paste0("model", i, ".csv"))  
+ }
```

Loops

- Example of a basic population projection based on a simple model:

$$p_{t+1} = p_t(1 + r_t)$$

```
> # initial population of 100 and growth rate of five percent per year
> p <- rep(NA, times = 10)
> p[1] <- 100
> for(i in 1:9){
+   p[i+1] <- p[i] * (1 + 0.05)
+ }
>
> p
[1] 100.0000 105.0000 110.2500 115.7625 121.5506 127.6282 134.0096 140.7100
[9] 147.7455 155.1328
```

Functions

```
> # can wrap for loops in a function
> sp_rate <- function(p0 = NULL, r = NULL, n = NULL){
+   p <- p0
+   for(i in 1:n){
+     p[i+1] <- p[i] * (1 + r)
+   }
+   return(p)
+ }
>
> sp_rate(p0 = 100, r = 0.05, n = 5)
[1] 100.0000 105.0000 110.2500 115.7625 121.5506 127.6282
>
> sp_rate(p0 = 20, r = 0.1, n = 20)
[1] 20.00000 22.00000 24.20000 26.62000 29.28200 32.21020 35.43122
[8] 38.97434 42.87178 47.15895 51.87485 57.06233 62.76857 69.04542
[15] 75.94997 83.54496 91.89946 101.08941 111.19835 122.31818 134.55000
```

Exercise 8 (ex18.R)

```
# 1. Create a function called tfr with inputs
#   a) asfr with NULL default
#   b) period with 5 as the default value
#   that calculates the total fertility rate (tfr) based on a vector
#   of age specific fertility rates (asfr) over a give period.
```

```
tfr <- function(##### = NULL, period = #####){
  x <- period * sum(asfr)
  return(x)
}
```

```
# 2. Run your function with the two sets of asfr values
#   a) 0.05, 0.1, 0.12, 0.09, 0.06 and 0.04
#   b) 0.1, 0.18, 0.16, 0.14, 0.12 and 0.02
```

```
tfr(asfr = c(0.05, #####, 0.12, 0.09, 0.06, #####))
#####(asfr = c(0.1, 0.18, 0.16, 0.14, #####, 0.02))
```

```
# 3. Create a function called rpois_add with inputs
#   a) n with default value 10
#   b) lambda1
#   c) lambda2
```

```
#   that adds together n random numbers from Poisson distributions with inputs from
```

```
rpois_add <- #####(n = 10, ##### = NULL, lambda2 = NULL){
  ##### <- rpois(n = #####, lambda = lambda1)
  r2 <- rpois(n = n, ##### = lambda2)
  #####(r1 + #####)
}
```

```
# 4. Run your function with n = 20, lambda1 = 5, lambda2 = 8
```


Working Directory

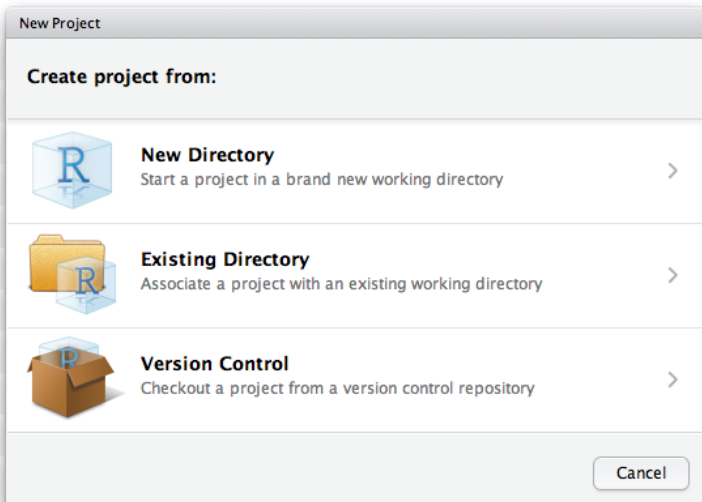
- R works in a particular location on your computer.
- The `getwd()` tells you the current R is working directory

```
> getwd()
"C:/Users/Guy/Teaching"
```

- We can change the location to make it easier to
 - Read in data (more on this later)
 - Run R scripts
 - Save data, plots, other outputs (more on this later).

Working Directory

- RStudio project files tells R which folder to work in.
 - Load the project file to get R set to the directory that you want to work in.



RStudio Projects

- At the start of your R session use File | Open Project
 - Can also select from the Project drop down in the top right hand corner.
- Once loaded or created you R session will be in the location of the .Rproj file.
- Can load data, save plots or source R scripts using full or relative paths
- For full paths R uses / or \\ instead of \ when setting directories:

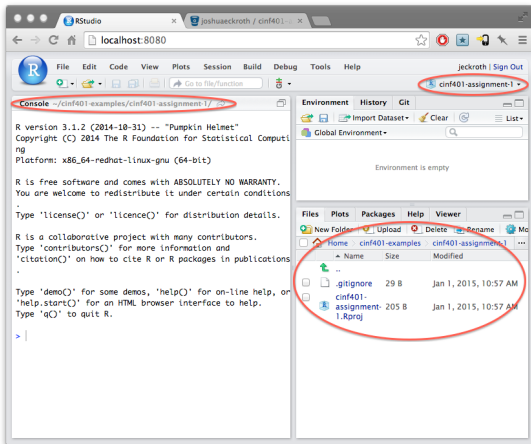
```
# read data from data folder
> d0 <- read_csv("C:/Users/Guy/Teaching/data/survey.csv")
# same as
> d0 <- read_csv("C:\\Users\\Guy\\Teaching\\data\\survey.csv")
# will not work
> d0 <- read_csv("C:\\Users\\Guy\\Teaching\\data\\survey.csv")
Error: '\\T' is an unrecognized escape in character string starting ""E:\\T"
```

Relative paths

- For relative paths
 - / is the root folder of the file system.
 - ./ usually denotes the current folder that your program or script is in, usually the same one with the file you run.
 - ../ denotes the folder above the current one.

```
# read from data folder in current directory - C:/Users/Guy/Teaching
> d0 <- read_csv("./data/survey.csv")
# read from another directory - C:/Users/Guy/Research
> d0 <- read_csv("../Research/data-raw/SurveyResults.csv")
```


RStudio Projects



Sourcing a Script

- The `source()` function runs R code in a file.

```
> # remove all object  
> rm(list = ls())  
> # run code in ex19.R using relative path:  
> source(file= "../exercise-solutions/ex18.R")
```



Sourcing a Script

- Using R code scripts (.R files) allows you to save and share your code.
 - Can break big project into multiple scripts.
 - Keeps a tidy work flow
 - Keeps the number of objects to a minimum
 - Quickly identify a particular line of code for revision.